

Homework 8 - Berkeley STAT 157

Prerequisites - Load Data

```
In [63]: import urllib3
import collections
import re
shakespeare = 'http://www.gutenberg.org/files/100/100-0.txt'

http = urllib3.PoolManager()
text = http.request('GET', shakespeare).data.decode('utf-8')
raw_dataset = ' '.join(re.sub('[^A-Za-z]+', ' ', text).lower().split())

print('number of characters: ', len(raw_dataset))
print(raw_dataset[0:70])
```

```
number of characters: 5032359
project gutenberg s the complete works of william shakespeare by willi
```

This dataset is quite a bit bigger than the time machine (5 million vs. 160k). For convenience we also include the remaining preprocessing steps. A bigger dataset will allow us to generate more meaningful models.

```
In [64]: idx_to_char = list(set(raw_dataset))
char_to_idx = dict([(char, i) for i, char in enumerate(idx_to_char)])
# vocabulary size, i.e., the number of letters in English
vocab_size = len(char_to_idx)
# whole corpus
corpus_indices = [char_to_idx[char] for char in raw_dataset]
sample = corpus_indices[:20]
print('chars:', ''.join([idx_to_char[idx] for idx in sample]))
print('indices:', sample)
# for training
train_indices = corpus_indices[:-100000]
# for testing
test_indices = corpus_indices[-100000:]
```

chars: project gutenber g

indices: [22, 19, 21, 16, 14, 18, 9, 13, 7, 0, 9, 14, 26, 5, 14, 19, 7, 13, 1, 13]

Lastly we import other useful libraries to help you getting started.

```
In [65]: import d2l
import math
import mxnet as mx
from mxnet import autograd, gluon, init, nd
from mxnet.gluon import loss as gloss, nn, rnn
import time
```

1. Train Recurrent Latent Variable Models

Train a number of different latent variable models using `train_indices` to assess their performance. By default pick 256 dimensions for the hidden units. You can use the codes provided in the class. Also, we strongly encourage you to use the Gluon implementation since it's a lot faster than building it from scratch.

1. Train a single-layer RNN (with latent variables).
2. Train a single-layer GRU.
3. Train a single-layer LSTM.
4. Train a two-layer LSTM.

How low can you drive the perplexity? Can you reproduce some of Shakespeare's finest writing (generate 200 characters). Start the sequence generator with `But Brutus is an honorable man`. Experiment with a number of settings:

- Number of hidden units.
- Embedding length.
- Gradient clipping.
- Number of iterations.
- Learning rate.

Save the models (at least in memory since you'll need them in the next exercise).

```
In [66]: def predict_rnn_gluon(prefix, num_chars, model, vocab_size, ctx, idx_to_char, char_to_idx):
    state = model.begin_state(batch_size=1, ctx=ctx)
    output = [char_to_idx[prefix[0]]]
    for t in range(num_chars + len(prefix) - 1):
        X = nd.array([output[-1]], ctx=ctx).reshape((1, 1))
        (Y, state) = model(X, state)
        if t < len(prefix) - 1:
            output.append(char_to_idx[prefix[t + 1]])
        else:
            output.append(int(Y.argmax(axis=1).asscalar()))
    return ''.join([idx_to_char[i] for i in output])

num_hiddens = 256
rnn_layer = rnn.RNN(num_hiddens)
rnn_layer.initialize()
num_steps = 35
num_epochs, batch_size, lr, clipping_theta = 200, 32, 30, 1e-2
```

```
In [67]: ctx = d2l.try_gpu()
single_rnn_model = d2l.RNNModel(rnn_layer, vocab_size)
single_rnn_model.initialize(force_reinit=True, ctx = ctx)
pred_period, pred_len, prefixes = 25, 100, ['but brutus is an honorable man']
```

Question 1 General Discoveries

- Learning rate: the model with smaller learning rate shows better performance. $lr = 30$ in our code.
- Hidden Units: the model with larger hidden units shows better performance, but it could not be too large. *hidden_units* = 320 in the GRU model.

- Embedding length: It seems that more we embedded, more precise the model would be. However, the memory of the computer need to memorize more. So we keep it as the same.
- Gradient clipping: We tried $1e-3$ and $1e-1$, the model shows worse performance. Therefore, we remian the clipping parameter unchanged.

Question 1.1: Single-layer RNN

[illegible]

epoch 25, perplexity 3.698738, time 19.11 sec

- but brutus is an honorable man s sooth the state the state the state the state the state the state the state the state the state the state t

epoch 50, perplexity 3.644675, time 18.97 sec

- but brutus is an honorable man s son and the state and the state and the state and the state and the state and the state and the s

epoch 75, perplexity 3.621387, time 18.90 sec

- but brutus is an honorable man s soul that i shall be the state the state the state the state the state the state the s
tate the st

epoch 100, perplexity 3.606846, time 18.64 sec

- but brutus is an honorable man and the service and the service and the service and the service and the service and the service and

epoch 125, perplexity 3.597160, time 18.68 sec

- but brutus is an honorable man s soul and the state they shall be the state they shall be the state they shall be the s
tate they s

epoch 150, perplexity 3.590280, time 19.11 sec

- but brutus is an honorable man s soul and the sea and the sea and the sea and the sea and the sea and the sea and the sea and the sea and the

epoch 175, perplexity 3.584725, time 18.54 sec

- but brutus is an honorable man s soul and the sea and the sea and the sea and the sea and the sea and the sea and the sea and the sea and the

epoch 200, perplexity 3.580413, time 18.83 sec

- but brutus is an honorable man s son and the sea and the sea and the sea and the sea and the sea and the sea and the se
a and the s


```
In [73]: predict_rnn_gluon('but brutus is an honorable man', 200, single_gru_model, vocab_size, ctx,\
                           idx_to_char, char_to_idx)
```

```
Out[73]: 'but brutus is an honorable man s and the sea and the country s son and then they say the stage and there s a present and
so i will not see thee when i shall see thee when i shall see thee when i shall see thee when i shall see the'
```

Question 1.4: Double Layer LSTM -- Best Performance

```
In [74]: # for question 1.4, double layers lstm
lstm_layer = rnn.LSTM(num_hiddens, num_layers=2)
double_lstm_model = d2l.RNNModel(lstm_layer, vocab_size)
d2l.train_and_predict_rnn_gluon(double_lstm_model, num_hiddens, vocab_size, ctx, train_indices, idx_to_char,\
                                char_to_idx, num_epochs, num_steps, lr, clipping_theta, batch_size,\
                                pred_period, pred_len, prefixes)
```

epoch 25, perplexity 3.214155, time 30.25 sec

- but brutus is an honorable man that s a foot of the country s head and then they say they say they say they say they sa
y they say

epoch 50, perplexity 3.007819, time 30.38 sec

- but brutus is an honorable man that would be so shall i say the countess is the sea and the senate have and they are th
ey are they

epoch 75, perplexity 2.915687, time 30.43 sec

- but brutus is an honorable man the thing is the fairest enter servant to the castle enter sir toby and sir toby and sir
toby and s

epoch 100, perplexity 2.861589, time 30.57 sec

- but brutus is an honorable man the third time the sea what s the matter from the capitol the commons have their shame t
hat they wo

epoch 125, perplexity 2.824575, time 30.29 sec

- but brutus is an honorable mantle s brother s life the first the best and so shall be the most unsure they would not ha
ve the sena

epoch 150, perplexity 2.799474, time 30.15 sec

- but brutus is an honorable man that s the song the fight of heaven and the son of my soul the streets of engling hands
and their t

epoch 175, perplexity 2.780415, time 30.20 sec

- but brutus is an honorable man may well be so betime as i am a soldier that i have seen the day will be the search d he
ad of the c

epoch 200, perplexity 2.764609, time 31.21 sec

- but brutus is an honorable man s house enter provost and so please you sir toby what she says she did and then to see t
he stars of

```
In [75]: predict_rnn_gluon('but brutus is an honorable man', 200, double_lstm_model, vocab_size, ctx,\n                        idx_to_char, char_to_idx)
```

```
Out[75]: 'but brutus is an honorable man s house enter provost and so please you sir toby what she says she did and then to see the\nstars of men the sea who should say they were not a stranger to the senate house s palace stephano i say sir'
```

2. Test Error

So far we measured perplexity only on the training set.

1. Implement a perplexity calculator that does not involve training.
2. Compute the perplexity of the best models in each of the 4 categories on the test set. By how much does it differ?

```
In [76]: # for question 2
# This function is used to test perplexity on test set, not including training
def test_perplexity(model, ctx, corpus_indices, idx_to_char, char_to_idx, num_epochs, num_steps, batch_size):
    for epoch in range(num_epochs):
        loss = gloss.SoftmaxCrossEntropyLoss()
        l_sum, n, start = 0.0, 0, time.time()
        data_iter = d2l.data_iter_consecutive(
            corpus_indices, batch_size, num_steps, ctx)
        state = model.begin_state(batch_size=batch_size, ctx=ctx)
        for X, Y in data_iter:
            for s in state:
                s.detach()
            with autograd.record():
                (output, state) = model(X, state)
                y = Y.T.reshape((-1,))
                l = loss(output, y).mean()
            l_sum += l.asscalar() * y.size
            n += y.size
        print('epoch %d, perplexity %f, time %.2f sec' % (
            epoch + 1, math.exp(l_sum / n), time.time() - start))
```

```
In [77]: num_epochs, num_steps, batch_size = 5, 50, 1
```



```
In [82]: # test perplexity single rnn
test_perplexity(single_rnn_model, ctx, test_indices, idx_to_char, char_to_idx, num_epochs, num_steps, batch_size)
```

```
epoch 1, perplexity 4.705128, time 3.91 sec
epoch 2, perplexity 4.706620, time 3.42 sec
epoch 3, perplexity 4.704766, time 3.39 sec
epoch 4, perplexity 4.704266, time 3.37 sec
epoch 5, perplexity 4.704968, time 3.38 sec
```

```
In [79]: # test perplexity on single lstm
test_perplexity(single_lstm_model, ctx, test_indices, idx_to_char, char_to_idx, num_epochs, num_steps, batch_size)
```

```
epoch 1, perplexity 4.527095, time 3.33 sec
epoch 2, perplexity 4.526816, time 3.42 sec
epoch 3, perplexity 4.527897, time 3.64 sec
epoch 4, perplexity 4.525007, time 3.63 sec
epoch 5, perplexity 4.524860, time 3.47 sec
```

```
In [80]: # test perplexity on gru
test_perplexity(single_gru_model, ctx, test_indices, idx_to_char, char_to_idx, num_epochs, num_steps, batch_size)
```

```
epoch 1, perplexity 4.599602, time 3.33 sec
epoch 2, perplexity 4.598010, time 3.42 sec
epoch 3, perplexity 4.597947, time 3.44 sec
epoch 4, perplexity 4.597218, time 3.42 sec
epoch 5, perplexity 4.597271, time 3.39 sec
```

```
In [81]: # test perplexity on double lstm
test_perplexity(double_lstm_model, ctx, test_indices, idx_to_char, char_to_idx, num_epochs, num_steps, batch_size)
```

```
epoch 1, perplexity 5.077668, time 7.28 sec
epoch 2, perplexity 5.082510, time 4.92 sec
epoch 3, perplexity 5.080637, time 4.86 sec
epoch 4, perplexity 5.082147, time 5.42 sec
epoch 5, perplexity 5.079242, time 5.48 sec
```

```
In [ ]:
```

